Alternate Model Comparison

Matched Pilot Results

English WikiANN named-entity recognition pilot. All three runs used the same fixed sample and recipe.

Model	Precision	Recall	F1	Validation loss
ModernBERT-large	0.301	0.256	**0.277**	1.206
Ettin-encoder-400M	0.238	0.147	0.181	1.398
DeBERTa-v3-large	0.283	0.078	0.122	1.346

Run Settings

Setting	Value
Training sample	240 sentences
Validation sample	120 sentences
Epochs	1
Batch size	8
Maximum sequence length	128 tokens
Learning rate	1e-5
Random seed	20261022
Weight precision	float32

Interpretation

ModernBERT had the highest F1 and recall in this pilot. DeBERTa's precision was close to ModernBERT's, but its recall was substantially lower. All validation losses were finite after loading every backbone in float32.

This is a small general-NER pilot, not a legal extraction evaluation or a reliable model ranking. The dataset is Wikipedia text, the sample is small, and training lasted one epoch. These metrics do not establish performance on Likewise fields, legal prose, negation, attribution, hedging, assertion precision, or coverage. Do not use this table to select a production extractor. A decision requires adjudicated legal-prose spans and held-out legal guard cases.

The dataset card does not declare a license. Keep this pilot report local and do not redistribute the dataset or represent these scores as an official benchmark.